


# DATA OF FUTURE PAST

## POSTGRES AS DISTRIBUTED ONLINE PROCESSING ANALYTICS ENGINE



by Gavin McQuillan / @gmcquillan

# SETTING

Data Engineering at Urban Airship, a mobile messaging company:

- Counting lots of things as fast as possible
- HBase to the rescue
- Home grown dimensional storage called **datacube**

# OUR TOPIC

1. Problem Statement
2. Exploration of Solutions
3. Benchmarking Solutions

# THE PROBLEM

- Data consistency
- New dimensions multiply writes
- Double counting
- Changing schema is hard
- Consistent backups?

# EXPLORING SOLUTIONS

Postgres is pretty nice to work with.

Makes adhoc analytics simple.

Well known replication and backup story

# PROBLEMS WITH POSTGRES

Not particularly good at scaling writes horizontally

Operationally complex

# EXISTING SOLUTIONS

- Postgres-xc/xl
- Slony
- Redshift
- Pg\_shard
- PLProxy

# PLPROXY

- Simple API
- Battle tested
- Flexible
- Easy upgrade paths, no lock-in



# APPROACH

Two phase commit

Commutative, Idempotent data

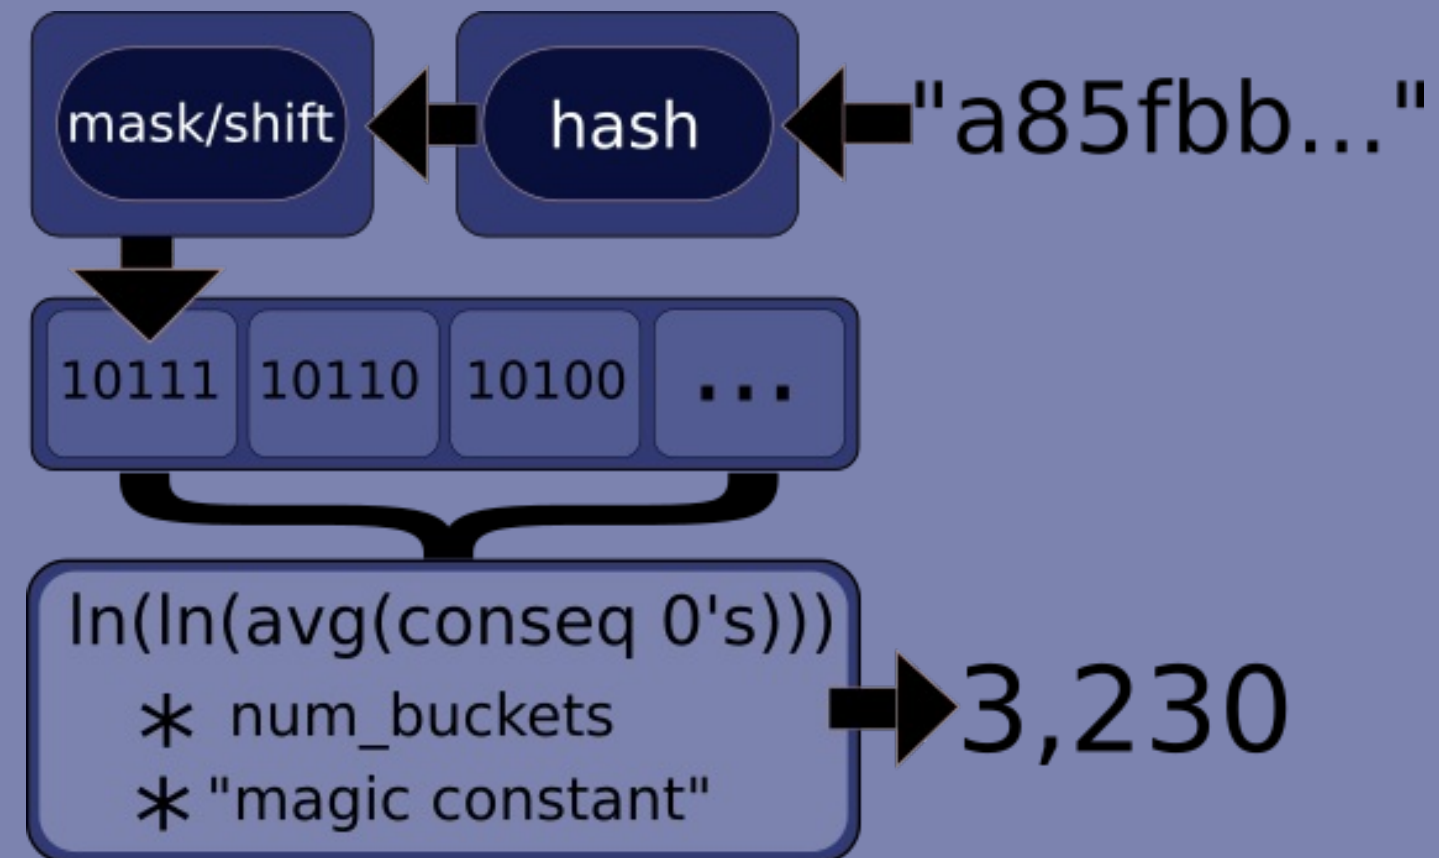
# IDEMPOTENT WRITES WITH HYPERLOGLOG

Postgres-hll extension

Commutative, idempotent

Fast, approximate, cardinality

# BRIEFLY, HOW HYPERLOGLOG WORKS



# PLPROXY: SETTING UP FOREIGN DATA WRAPPERS IN SQL.

# CLUSTER CONFIG

Partition defs, cluster version, connection config elided

Partition mapping is as follows:

# PARTITION MAPPING

```
CREATE FOREIGN DATA WRAPPER plproxy;

CREATE SERVER naboocounts FOREIGN DATA WRAPPER plproxy
OPTIONS (connection_lifetime '1800',
         p0 'dbname=part00 host=10.130.1.38',
         p1 'dbname=part01 host=10.130.1.39' );

-- This mapping is accessible to all local users
CREATE USER MAPPING FOR PUBLIC SERVER naboocounts;
```

# PROXY FUNCTIONS

```
CREATE OR REPLACE FUNCTION upsert_count(  
    in_id text, in_date date, in_hour smallint,  
    in_event_id text, in_category text  
) RETURNS TABLE (updates int)  
    LANGUAGE plproxy  
    AS $$  
    CLUSTER 'testcounts';  
    SPLIT ALL;  
    RUN ON hashtext(in_event_id);  
    $$;
```

# EXPERIMENTAL DESIGN

Can a sharded postgres reasonably keep up with Hbase?




# PHYSICAL LAYOUT

Three Dell R610s with:

- 2 8-core Xeon CPUs
- 6 SSDs in a RAID 10 configuration (~300GB usable)
- write-back cache enabled on the I/O controller
- 48GB of ECC RAM.
- Bonded Ethernet interfaces

# SIMPLE TOPOLOGY



# SETTING UP THE SHARDS

# EXAMPLE TABLE

```
CREATE TABLE test_counts
(
    id CHAR(22),
    date DATE,
    hour SMALLINT,
    event_ids hll,
    category TEXT
);
```

# SINGLE INSERT/UPDATE

```
CREATE OR REPLACE FUNCTION upsert_test_count(...) RETURNS int
BEGIN
    UPDATE test_counts set event_ids=hll_add(
        event_ids, hll_hash_text(in_event_id))
        WHERE ...
    IF FOUND THEN RETURN 0; END IF;
    BEGIN
        INSERT INTO test_counts(event_ids, ...)
            VALUES (hll_empty(), ...);
        Update test_counts SET event_ids=hll_add(
            event_ids, hll_hash_text(in_event_id))
            WHERE ...
    END;
    RETURN 1;
END;
```

# SINGLE WRITE

```
select upsert_test_count(  
    'some-identifier-string'::text,  
    '2015-05-16'::date,  
    '22'::smallint,  
    'cabef32d-bc21-4a34-993d-3e7d606df9c6'::text,  
    'Catagory1'::text  
);
```

# INDEXES

Optimum index configuration (3/4 dimensions indexed):

```
create index on test_counts (id, date, hour) with(fillfactor=10);
```

The fillfactor tells Postgres to pre-allocate 90% of the index space empty, copy data less.

# POSTGRES TUNING PARAMETERS

Config Parameter	My Setting	Notes
shared_buffers	10GB	25% of memory
effective_cache_size	30GB	75% of memory
work_mem	50MB	sorting, etc. it's also per-session
random_page_cost	2.0	default is 4.0; SSDs here.
checkpoint_segments	256	default 32; number WAL segments to buffer



# STILL TOO SLOW

~2,000 events/sec

A transaction per tuple just won't work long-term


# BATCHING

```
CREATE OR REPLACE FUNCTION upsert_test_count(  
    in_ids text[], in_dates date[], in_hours smallint[],  
    in_event_ids text[], in_cats text[]  
) RETURNS TABLE (update int)  
BEGIN  
  
RETURN QUERY SELECT upsert_push_hll(  
    c.in_ids, c.in_date, c.in_hour, c.in_event_id, c.in_cats  
) FROM unnest(  
    in_ids, in_dates, in_hours, in_event_ids, in_cats  
) as c (in_id, in_date, in_hour, in_event_id, in_cats);  
END;  
$;
```


# BATCH WRITE QUERY

```
select upsert_test_count(  
    '{aaaaaaaaaaaaaaaaaaaaaaa, ..., ...}'::text[],  
    '{2015-05-15,2015-05-16,2015-05-16}'::date[],  
    '{20,21,23}'::smallint[],  
    '{cabef32d-bc21-4a34-993d-3e7d606df9b1, ..., ...}'::text[],  
    '{Category1,Catagory2,Category1}'::text[]  
);
```


# ANATOMY OF A PLPROXY TRANSACTION



# WHEN THINGS GO WRONG



# DEADLOCK DETECTED!



# DEADLOCK SOLUTIONS

- Sort tuples before submitting them
- Single writer pattern

Our functions make sorting difficult, so single writer

# SIMPLE TOPOLOGY


Peaks out with tuning, indexes, and batching at **11k events/sec**

Next step is to increase parallelism




# BENCHMARK RESULTS


# SIMPLE TOPOLOGY THROUGHPUT (200K)



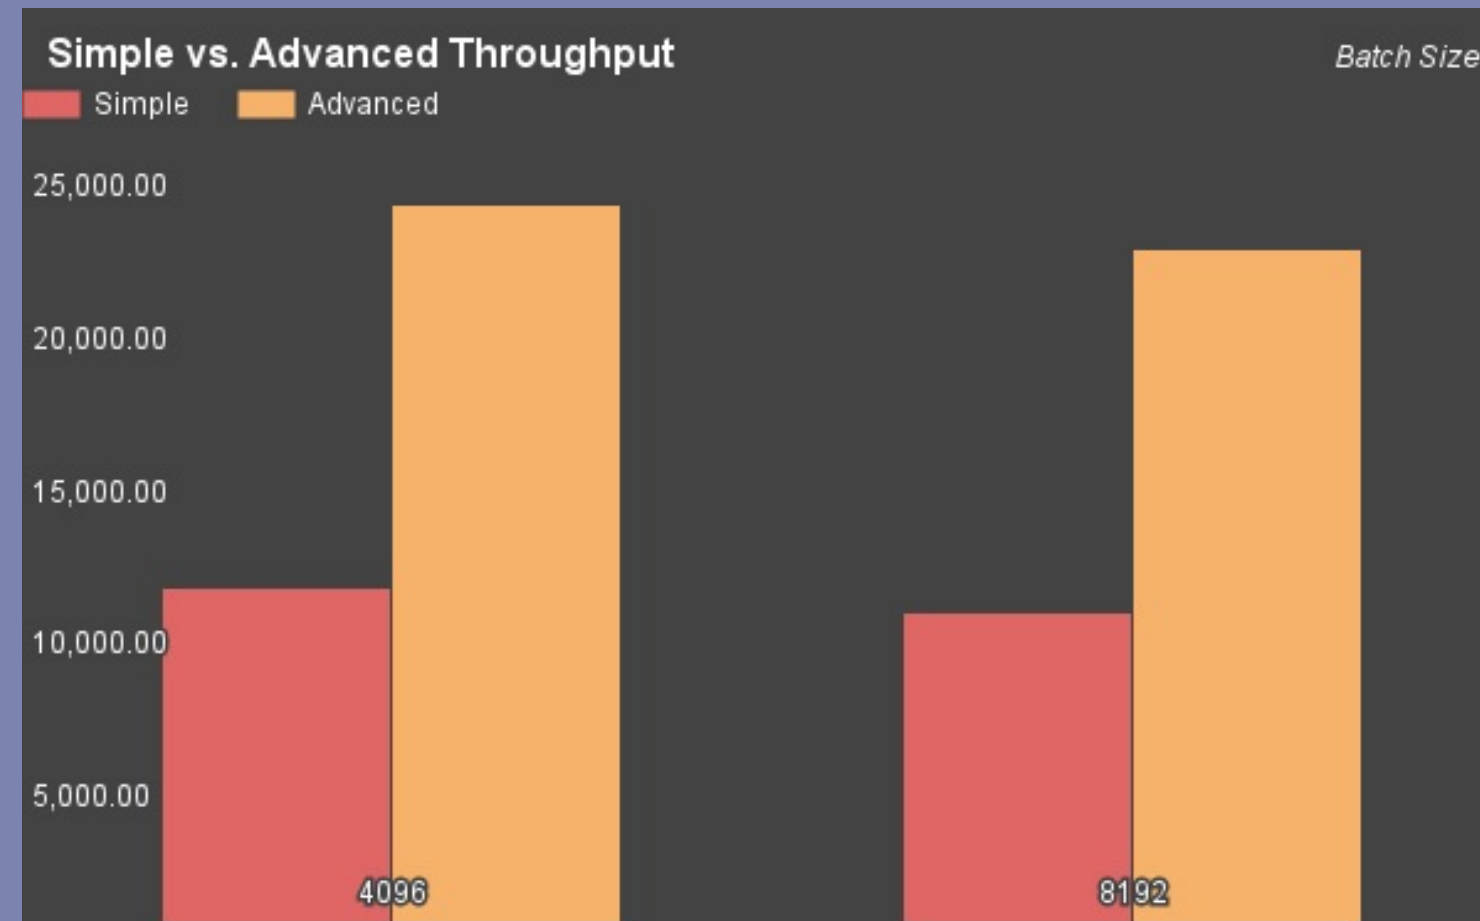
# ADVANCED TOPOLOGY



# ADVANCED TOPOLOGY THROUGHPUT (2MM)



# DIRECT COMPARISON (2MM)



**LET'S SEE RESULTS ON  
A LOADED CLUSTER!**

# TYPES OF LOAD

1. Data load: number of rows, size on disk
2. Concurrent requests

# SETTING UP A LOADED SYSTEM

Cluster State	Cluster Size (MB)	Index Size(MB)	Number of Rows	Number of IDs
Before	63,864	21,824	333,757,839	307,520
After	80,096	27,088	412,900,728	357,520




# SETTING UP CONCURRENT REQUESTS

Pre-generate insert query batches into .sql files

Run 10 concurrently in a screen session

Not 100% representative of real-world behavior

# LOADED RESULTS



# READ QUERY

```
SELECT id, date, hour, hll_cardinality(event_ids)
FROM dynamic_query(
    'SELECT * from test_counts
    WHERE date >= (now() - interval ''7 days'')
    AND id = ''M2E0MDd1NzYtY2Y4NC00Nz'''
) AS (
    id char(22), date date, hour smallint, event_ids hll, cat text)
ORDER BY
    date desc,
    hour desc
LIMIT 10;
```

# READ QUERY RESULTS

id		date		hour		hll_cardinality
-----+-----+-----						
M2E0MDd1NzYtY2Y4NC00Nz		2015-06-10		18		6
M2E0MDd1NzYtY2Y4NC00Nz		2015-06-10		13		6
M2E0MDd1NzYtY2Y4NC00Nz		2015-06-10		13		6
M2E0MDd1NzYtY2Y4NC00Nz		2015-06-10		6		6
M2E0MDd1NzYtY2Y4NC00Nz		2015-06-10		21		5

# READ LATENCY

Generally 10's to 100's of ms

# WRAP UP: POSTGRES FOR DISTRIBUTED OLAP

- Postgres can scale horizontally.
- Write throughput  $\approx$  Hbase system.
- New features are a few lines of SQL
- We **retain** queryability and DDLs
- Operational concerns only get worse :(

# REMAINING WORK

# FUTURE FEATURES

- Cross table joins
- Automated failovers(shards)
- Automated, efficient backups
- Tools to help migrate data, add partitions
- Integrating PGBouncer



# WORK IS ONGOING

Ansible automation for setting up a test cluster

[github.com/gmcquillan/pg\\_plural](https://github.com/gmcquillan/pg_plural)

**THANK YOU**

# REFERENCES

- PLProxy Syntax Reference
- PLProxy FAQ
- Martin Kleppmann on Transactions [VIDEO]
- depesz.com
- Urbanski Presentation at pgconf.ru [PDF]
- Deadlocks in Postgresql
- HyperLogLog: the analysis of near-optimal cardinality estimation algorithm - Flajolet [PDF]